

Assignment - Decision Trees and Random Forests



In this assignment, you'll continue building on the previous assignment to predict the price of a house using information like its location, area, no. of rooms etc. You'll use the dataset from the [House Prices - Advanced Regression Techniques](#) competition on [Kaggle](#).

We'll follow a step-by-step process:

1. Download and prepare the dataset for training
2. Train, evaluate and interpret a decision tree
3. Train, evaluate and interpret a random forest
4. Tune hyperparameters to improve the model
5. Make predictions and save the model

As you go through this notebook, you will find a ??? in certain places. Your job is to replace the ??? with appropriate code or values, to ensure that the notebook runs properly end-to-end and your machine learning model is trained properly without errors.

Guidelines

1. Make sure to run all the code cells in order. Otherwise, you may get errors like `NameError` for undefined variables.
2. Do not change variable names, delete cells, or disturb other existing code. It may cause problems during evaluation.
3. In some cases, you may need to add some code cells or new statements before or after the line of code containing the ???.
4. Since you'll be using a temporary online service for code execution, save your work by running `jovian.commit` at regular intervals.
5. Review the "Evaluation Criteria" for the assignment carefully and make sure your submission meets all the criteria.
6. Questions marked (**Optional**) will not be considered for evaluation and can be skipped. They are for your learning.
7. It's okay to ask for help & discuss ideas on the [community forum](#), but please don't post full working code, to give everyone an opportunity to solve the assignment on their own.

Important Links:

- Make a submission here: <https://jovian.ai/learn/machine-learning-with-python-zero-to-gbms/assignment/assignment-2-decision-trees-and-random-forests>

- Ask questions, discuss ideas and get help here: <https://jovian.ai/forum/c/zero-to-gbms/gbms-assignment-2/99>
- Review this Jupyter notebook: <https://jovian.ai/aakashns/sklearn-decision-trees-random-forests>

How to Run the Code and Save Your Work

Option 1: Running using free online resources (1-click, recommended): The easiest way to start executing the code is to click the **Run** button at the top of this page and select **Run on Binder**. This will set up a cloud-based Jupyter notebook server and allow you to modify/execute the code.

Option 2: Running on your computer locally: To run the code on your computer locally, you'll need to set up [Python](#), download the notebook and install the required libraries. Click the **Run** button at the top of this page, select the **Run Locally** option, and follow the instructions.

Saving your work: You can save a snapshot of the assignment to your [Jovian](#) profile, so that you can access it later and continue your work. Keep saving your work by running `jovian.commit` from time to time.

```
!pip install jovian --upgrade --quiet
```

```
import jovian
```

```
jovian.commit(project='python-random-forests-assignment', privacy='secret')
```

```
[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on  
https://jovian.ai
```

```
[jovian] Committed successfully! https://jovian.ai/ch20b112/python-random-forests-assignment
```

```
'https://jovian.ai/ch20b112/python-random-forests-assignment'
```

Let's begin by installing the required libraries.

```
!pip install opendatasets scikit-learn plotly folium --upgrade --quiet
```

```
!pip install pandas numpy matplotlib seaborn xgboost --quiet
```

Download and prepare the dataset for training

```
import os  
from zipfile import ZipFile  
from urllib.request import urlretrieve  
  
dataset_url = 'https://github.com/JovianML/opendatasets/raw/master/data/house-prices-ac  
urlretrieve(dataset_url, 'house-prices.zip')  
with ZipFile('house-prices.zip') as f:  
    f.extractall(path='house-prices')
```

```
os.listdir('house-prices')
```

```
['data_description.txt', 'sample_submission.csv', 'test.csv', 'train.csv']
```

```
import pandas as pd
pd.options.display.max_columns = 200
pd.options.display.max_rows = 200

prices_df = pd.read_csv('house-prices/train.csv')
prices_df
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Inside
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	Front
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corn
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	Front
...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	Lvl	AllPub	Inside
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	Lvl	AllPub	Inside
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	Lvl	AllPub	Inside
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	Lvl	AllPub	Inside
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	Lvl	AllPub	Inside

1460 rows × 81 columns

```
import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.model_selection import train_test_split

# Identify input and target columns
input_cols, target_col = prices_df.columns[1:-1], prices_df.columns[-1]
inputs_df, targets = prices_df[input_cols].copy(), prices_df[target_col].copy()

# Identify numeric and categorical columns
numeric_cols = prices_df[input_cols].select_dtypes(include=np.number).columns.tolist()
categorical_cols = prices_df[input_cols].select_dtypes(include='object').columns.tolist()

# Impute and scale numeric columns
imputer = SimpleImputer().fit(inputs_df[numeric_cols])
inputs_df[numeric_cols] = imputer.transform(inputs_df[numeric_cols])
scaler = MinMaxScaler().fit(inputs_df[numeric_cols])
inputs_df[numeric_cols] = scaler.transform(inputs_df[numeric_cols])

# One-hot encode categorical columns
encoder = OneHotEncoder(sparse=False, handle_unknown='ignore').fit(inputs_df[categorical_cols])
```

```

encoded_cols = list(encoder.get_feature_names(categorical_cols))
inputs_df[encoded_cols] = encoder.transform(inputs_df[categorical_cols])

# Create training and validation sets
train_inputs, val_inputs, train_targets, val_targets = train_test_split(
    inputs_df[numeric_cols + encoded_cols], targets, test_size=0.25, random_state=42)

```

/opt/conda/lib/python3.9/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function get_feature_names is deprecated; get_feature_names is deprecated in 1.0 and will be removed in 1.2. Please use get_feature_names_out instead.

```
warnings.warn(msg, category=FutureWarning)
```

/opt/conda/lib/python3.9/site-packages/pandas/core/frame.py:3678: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe = frame.copy()`

```
self[col] = igetitem(value, i)
```

Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

```
inputs_df=inputs_df.drop(columns=categorical_cols)
```

```
inputs_df.isna().sum()
```

```

MSSubClass      0
LotFrontage     0
LotArea         0
OverallQual     0
OverallCond     0
..
SaleCondition_AdjLand  0
SaleCondition_Alloca  0
SaleCondition_Family  0
SaleCondition_Normal  0
SaleCondition_Partial  0
Length: 304, dtype: int64

```

Decision Tree

QUESTION 1: Train a decision tree regressor using the training set.

```
from sklearn.tree import DecisionTreeRegressor
```

```
# Create the model  
tree = DecisionTreeRegressor(random_state=42)
```

```
# Fit the model to the training data  
tree.fit(train_inputs, train_targets)
```

```
DecisionTreeRegressor(random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

DecisionTreeRegressor

```
DecisionTreeRegressor(random_state=42)
```

Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

QUESTION 2: Generate predictions on the training and validation sets using the trained decision tree, and compute the RMSE loss.

```
from sklearn.metrics import mean_squared_error
```

```
tree_train_preds = tree.predict(train_inputs)
```

```
tree_train_rmse = mean_squared_error(tree_train_preds, train_targets, squared=False)
```

```
tree_val_preds = tree.predict(val_inputs)
```

```
tree_val_rmse = mean_squared_error(tree_val_preds, val_targets, squared=False)
```

```
print('Train RMSE: {}, Validation RMSE: {}'.format(tree_train_rmse, tree_val_rmse))
```

Train RMSE: 0.0, Validation RMSE: 37802.872039112044

Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on
<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

QUESTION 3: Visualize the decision tree (graphically and textually) and display feature importances as a graph. Limit the maximum depth of graphical visualization to 3 levels.

```
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree, export_text
import seaborn as sns
sns.set_style('darkgrid')
%matplotlib inline
```

```
plt.figure(figsize=(30, 15))
```

```
# Visualize the tree graphically using plot_tree
plot_tree(tree, feature_names=inputs_df.columns, max_depth=3, filled=True);
```



```

|   |   |   |   |--- truncated branch of depth 9
|--- OverallQual > 0.72
|   |--- OverallQual <= 0.83
|   |   |--- GrLivArea <= 0.31
|   |   |   |--- BsmtFinSF1 <= 0.17
|   |   |   |   |--- truncated branch of depth 13
|   |   |   |   |--- BsmtFinSF1 > 0.17
|   |   |   |   |--- truncated branch of depth 12
|   |   |   |   |--- GrLivArea > 0.31
|   |   |   |   |--- 1stFlrSF <= 0.26
|   |   |   |   |--- truncated branch of depth 10
|   |   |   |   |--- 1stFlrSF > 0.26
|   |   |   |   |--- truncated branch of depth 7
|   |   |   |   |--- OverallQual > 0.83
|   |   |   |   |--- 2ndFlrSF <= 0.81
|   |   |   |   |--- LotFrontage <= 0.33
|   |   |   |   |--- truncated branch of depth 10
|   |   |   |   |--- LotFrontage > 0.33
|   |   |   |   |--- truncated branch of depth 3
|   |   |   |   |--- 2ndFlrSF > 0.81
|   |   |   |   |--- 2ndFlrSF <= 0.93
|   |   |   |   |--- value: [625000.00]
|   |   |   |   |--- 2ndFlrSF > 0.93
|   |   |   |   |--- value: [745000.00]

```

```

# Check feature importance
tree_importances = tree.feature_importances_

```

```

tree_importance_df = pd.DataFrame({
    'feature': train_inputs.columns,
    'importance': tree_importances
}).sort_values('importance', ascending=False)

```


```
tree_importance_df
```

	feature	importance
3	OverallQual	0.632537
15	GrLivArea	0.104334
13	2ndFlrSF	0.031896
11	TotalBsmtSF	0.028504
12	1stFlrSF	0.026730
...

	feature	importance
104	Condition2_RRAn	0.000000
103	Condition2_RRAe	0.000000
102	Condition2_PosN	0.000000
212	BsmtFinType2_nan	0.000000
152	Exterior2nd_CBlock	0.000000

304 rows × 2 columns

```
plt.title('Decision Tree Feature Importance')
sns.barplot(data=tree_importance_df.head(10), x='importance', y='feature');
```



Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on <https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

Random Forests

QUESTION 4: Train a random forest regressor using the training set.

```
from sklearn.ensemble import RandomForestRegressor
```

```
# Create the model
rf1 = RandomForestRegressor(random_state=42)
```

```
# Fit the model
rf1.fit(train_inputs, train_targets)
```

```
RandomForestRegressor(random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
RandomForestRegressor

```
RandomForestRegressor(random_state=42)
```

Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on
<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

QUESTION 5: Make predictions using the random forest regressor.

```
rf1_train_preds = rf1.predict(train_inputs)
```

```
rf1_train_rmse = mean_squared_error(rf1_train_preds, train_targets, squared=False)
```

```
rf1_val_preds = rf1.predict(val_inputs)
```

```
rf1_val_rmse = mean_squared_error(rf1_val_preds, val_targets, squared=False)
```

```
print('Train RMSE: {}, Validation RMSE: {}'.format(rf1_train_rmse, rf1_val_rmse))
```

Train RMSE: 11864.318299877767, Validation RMSE: 27830.03863639856

Let's save our work before continuing.

```
jovian.commit()
```

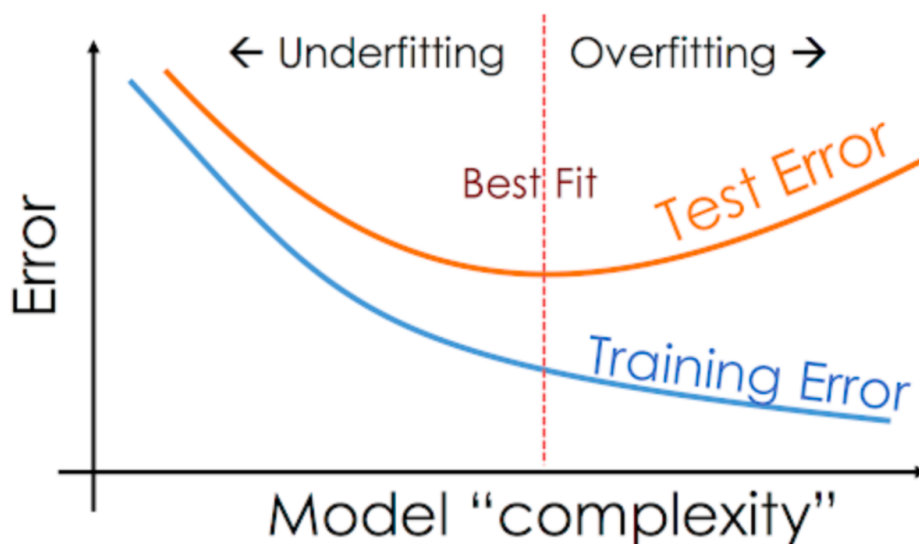
[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on <https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

Hyperparameter Tuning

Let us now tune the hyperparameters of our model. You can find the hyperparameters for RandomForestRegressor here: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>



Hyperparameters are use

Let's define a helper function `test_params` which can test the given value of one or more hyperparameters.

```
def test_params(**params):  
    model = RandomForestRegressor(random_state=42, n_jobs=-1, **params).fit(train_input  
    train_rmse = mean_squared_error(model.predict(train_inputs), train_targets, squared  
    val_rmse = mean_squared_error(model.predict(val_inputs), val_targets, squared=False  
    return train_rmse, val_rmse
```

It can be used as follows:

```
test_params(n_estimators=20, max_depth=20)
```

(13776.89957127333, 28886.033523273858)


```
test_params(n_estimators=50, max_depth=10, min_samples_leaf=4, max_features=0.4)
```

(20490.359632429263, 29804.931642791606)

Let's also define a helper function to test and plot different values of a single parameter.

```
def test_param_and_plot(param_name, param_values):  
    train_errors, val_errors = [], []  
    for value in param_values:  
        params = {param_name: value}  
        train_rmse, val_rmse = test_params(**params)  
        train_errors.append(train_rmse)  
        val_errors.append(val_rmse)  
    plt.figure(figsize=(10,6))  
    plt.title('Overfitting curve: ' + param_name)  
    plt.plot(param_values, train_errors, 'b-o')  
    plt.plot(param_values, val_errors, 'r-o')  
    plt.xlabel(param_name)  
    plt.ylabel('RMSE')  
    plt.legend(['Training', 'Validation'])
```

```
test_param_and_plot('max_depth', [5, 10, 15, 20, 25, 30, 35])
```



From the above graph, it appears that the best value for `max_depth` is around 20, beyond which the model starts to overfit.

Let's save our work before continuing.

```
jovian.commit()
```


[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on <https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'


QUESTION 6: Use the `test_params` and `test_param_and_plot` functions to experiment with different values of the hyperparameters like `n_estimators`, `max_depth`, `min_samples_split`, `min_samples_leaf`, `min_weight_fraction_leaf`, `max_features`, `max_leaf_nodes`, `min_impurity_decrease`, `min_impurity_split` etc. You can learn more about the hyperparameters here: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>

```
test_param_and_plot('n_estimators', [10,20,30,40,50,60,70,80,90,100])
```



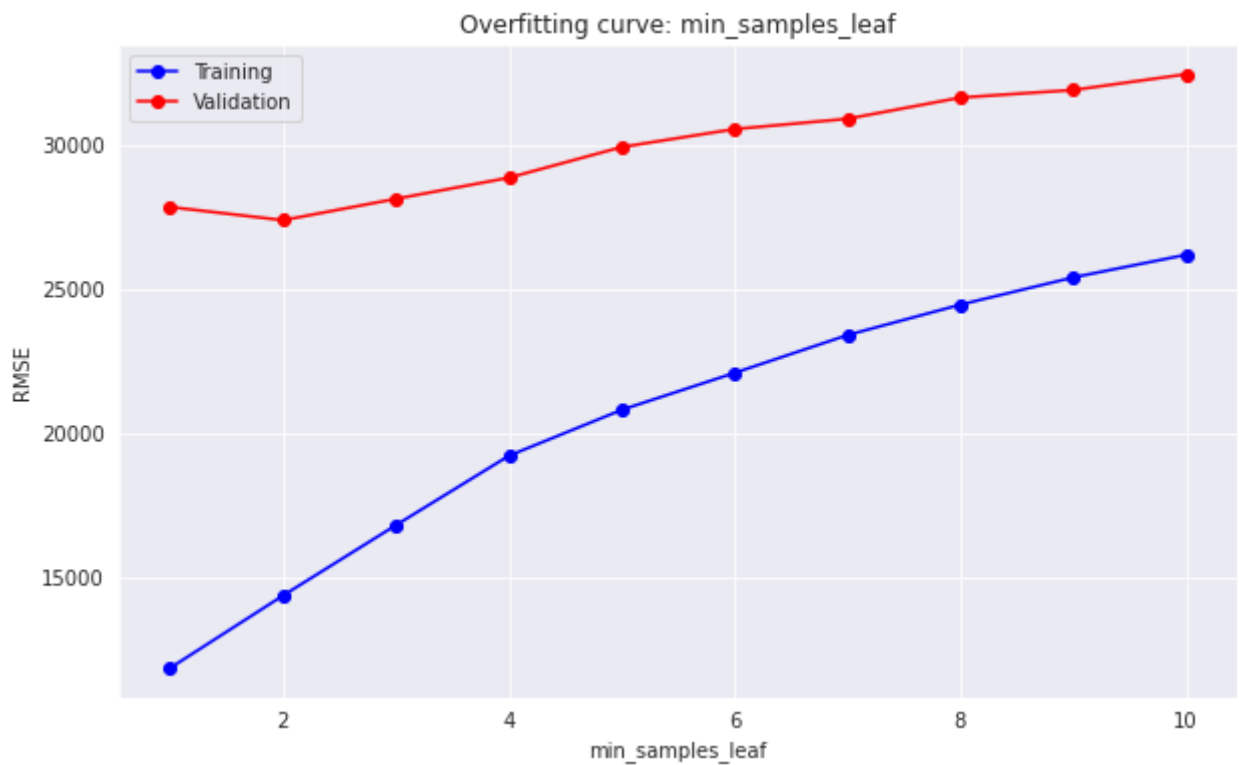
#from above graph it appears that best value of n_estimators=70

```
test_param_and_plot('min_samples_split', [2,3,4,5,6,7,8,9,10])
```



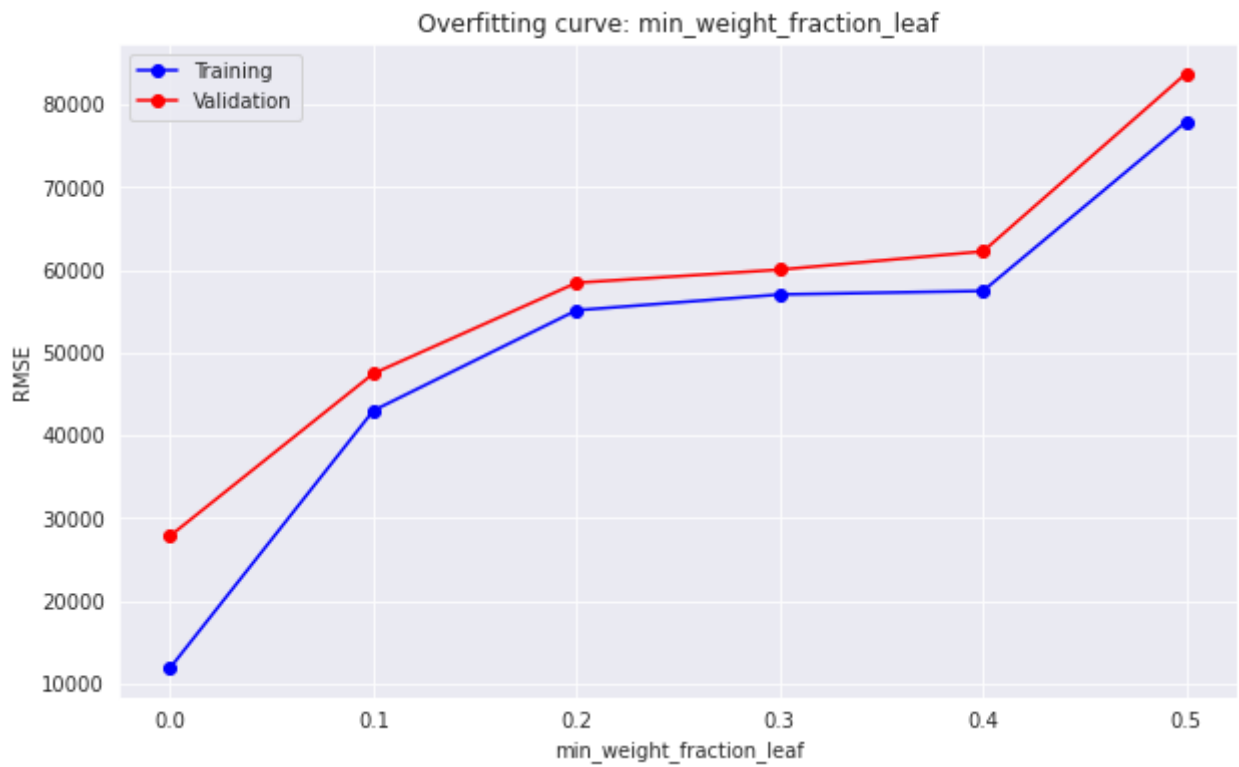
#from above graph it appears that best value of min_samples_split=7

```
test_param_and_plot('min_samples_leaf', [1,2,3,4,5,6,7,8,9,10])
```



#from above graph it appears that best value of min_samples_leaf=2

```
test_param_and_plot('min_weight_fraction_leaf', [0,0.1,0.2,0.3,0.4,0.5])
```



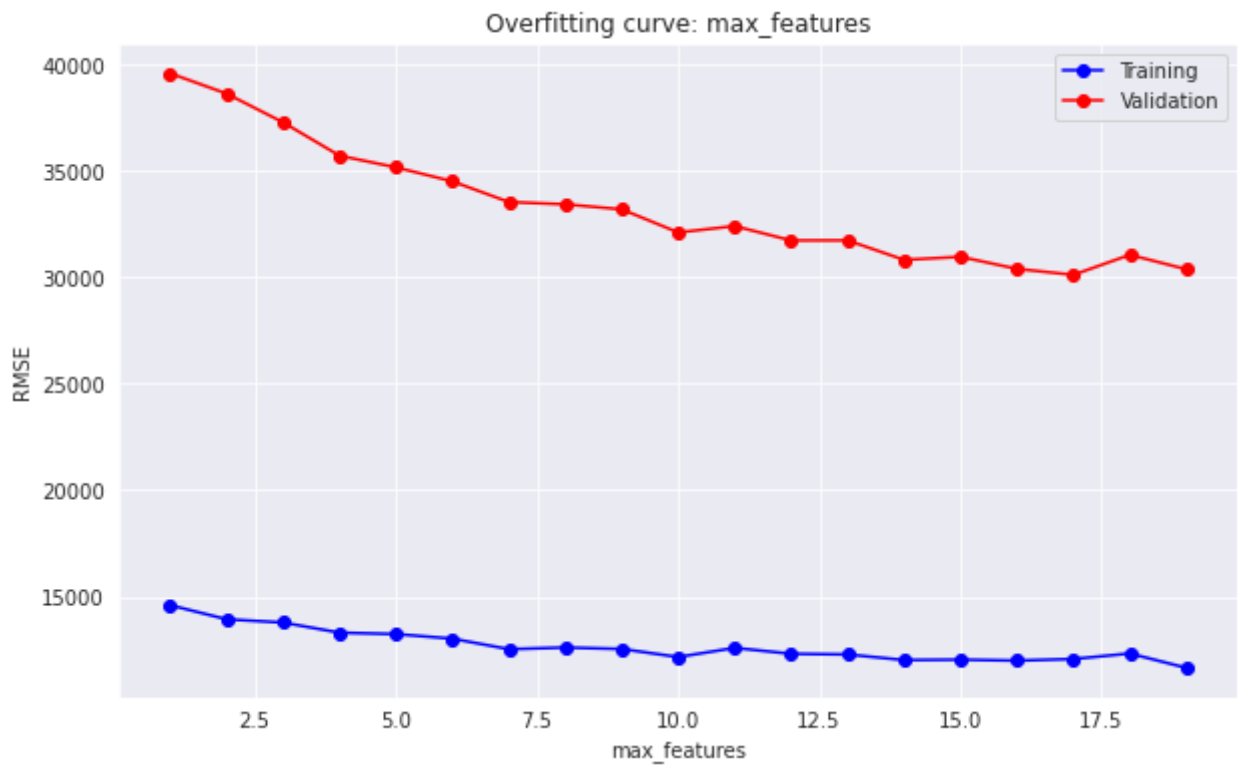
#from above graph it appears that best value of min_weight_fraction_leaf=0

inputs_df

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtF
0	0.235294	0.150685	0.033420	0.666667	0.500	0.949275	0.883333	0.122500	0.12
1	0.000000	0.202055	0.038795	0.555556	0.875	0.753623	0.433333	0.000000	0.17
2	0.235294	0.160959	0.046507	0.666667	0.500	0.934783	0.866667	0.101250	0.08
3	0.294118	0.133562	0.038561	0.666667	0.500	0.311594	0.333333	0.000000	0.03
4	0.235294	0.215753	0.060576	0.777778	0.500	0.927536	0.833333	0.218750	0.11
...
1455	0.235294	0.140411	0.030929	0.555556	0.500	0.920290	0.833333	0.000000	0.00
1456	0.000000	0.219178	0.055505	0.555556	0.625	0.768116	0.633333	0.074375	0.13
1457	0.294118	0.154110	0.036187	0.666667	1.000	0.500000	0.933333	0.000000	0.04
1458	0.000000	0.160959	0.039342	0.444444	0.625	0.565217	0.766667	0.000000	0.00
1459	0.000000	0.184932	0.040370	0.444444	0.625	0.673913	0.250000	0.000000	0.14

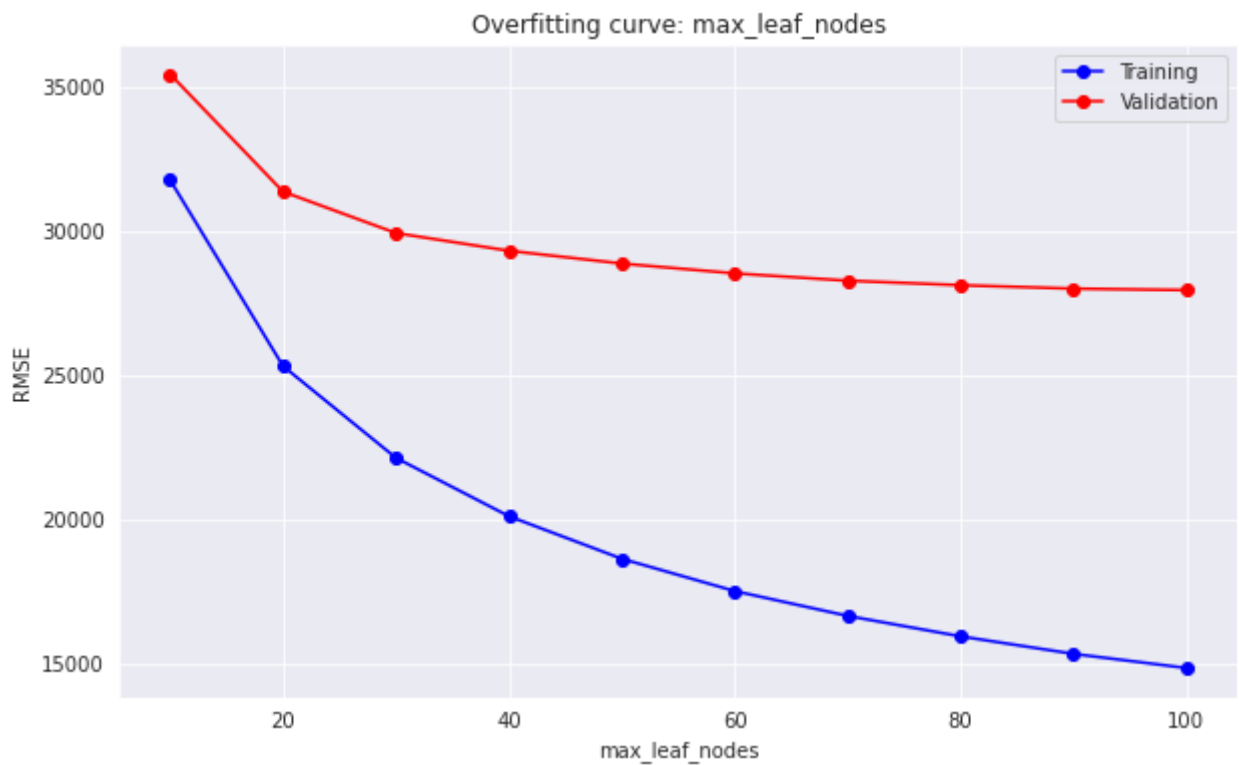
1460 rows × 304 columns

```
test_param_and_plot('max_features', list(range(1,20)))
```



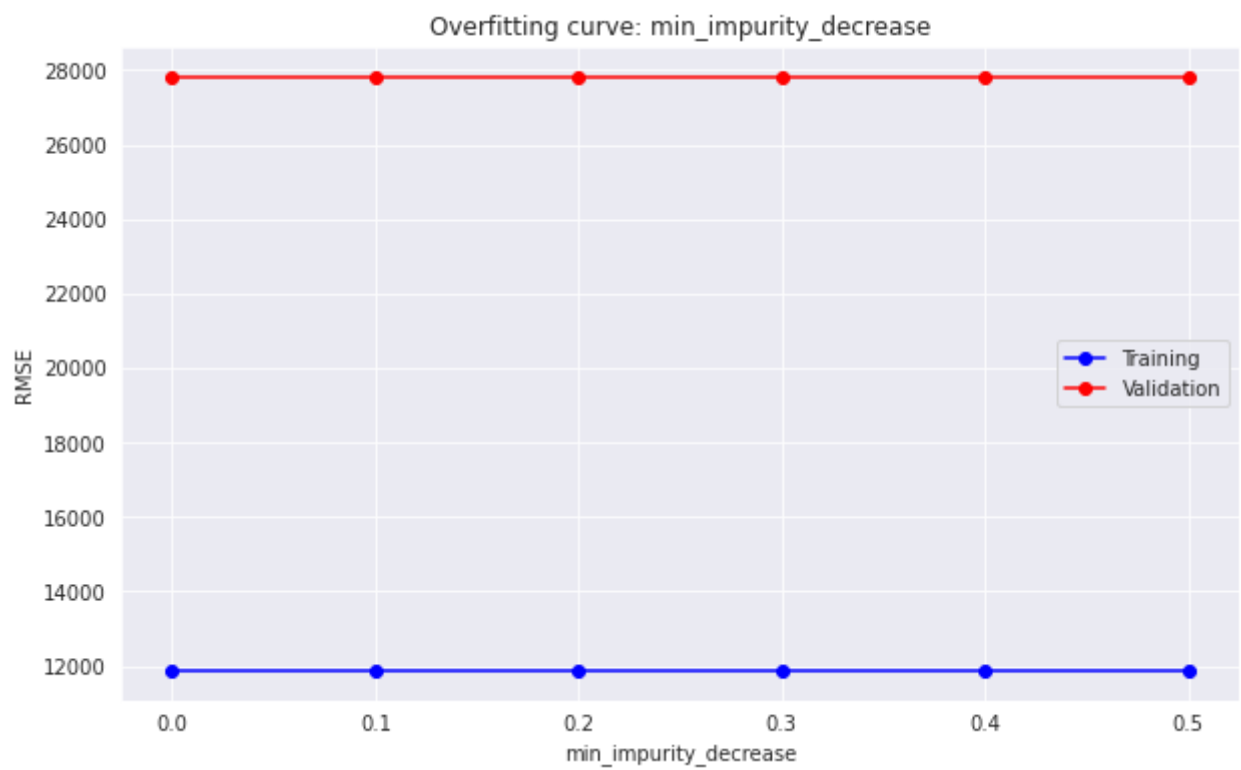
#from above graph it appears that best value of max_features=42

```
test_param_and_plot('max_leaf_nodes', [10,20,30,40,50,60,70,80,90,100])
```




#from above graph it appears that best value of max_leaf_nodes=40

```
test_param_and_plot('min_impurity_decrease', [0,0.1,0.2,0.3,0.4,0.5])
```

#from above graph it appears that best value of min_impurity_decrease=0

```
test_param_and_plot('max_samples', [100,200,300,400,500,600,700,800,900,1000])
```



Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on

<https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

Training the Best Model

QUESTION 7: Train a random forest regressor model with your best hyperparameters to minimize the validation loss.

```
# Create the model with custom hyperparameters
```

```
rf2 = RandomForestRegressor(random_state=42,max_depth=20,n_estimators=250,min_samples_s
```

```
# Train the model
```

```
rf2.fit(train_inputs,train_targets)
```

```
RandomForestRegressor(max_depth=20, max_features=200, max_leaf_nodes=100,  
min_impurity_decrease=0, min_samples_leaf=2,  
min_samples_split=3, min_weight_fraction_leaf=0,  
n_estimators=250, random_state=42)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
RandomForestRegressor

```
RandomForestRegressor(max_depth=20, max_features=200, max_leaf_nodes=100,  
min_impurity_decrease=0, min_samples_leaf=2,  
min_samples_split=3, min_weight_fraction_leaf=0,  
n_estimators=250, random_state=42)
```

```
from xgboost import XGBRegressor
```

```
xg=XGBRegressor(random_state=42,max_depth=50,n_estimators=250,learning_rate=0.06)  
xg.fit(train_inputs,train_targets)
```

```
XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,  
colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,  
early_stopping_rounds=None, enable_categorical=False,  
eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
```

```
importance_type=None, interaction_constraints='',
learning_rate=0.06, max_bin=256, max_cat_to_onehot=4,
max_delta_step=0, max_depth=50, max_leaves=0, min_child_weight=1,
missing=nan, monotone_constraints='()', n_estimators=250, n_jobs=0,
num_parallel_tree=1, predictor='auto', random_state=42,
reg_alpha=0, reg_lambda=1, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.
XGBRegressor

```
XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
early_stopping_rounds=None, enable_categorical=False,
eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
importance_type=None, interaction_constraints='',
learning_rate=0.06, max_bin=256, max_cat_to_onehot=4,
max_delta_step=0, max_depth=50, max_leaves=0, min_child_weight=1,
missing=nan, monotone_constraints='()', n_estimators=250, n_jobs=0,
num_parallel_tree=1, predictor='auto', random_state=42,
reg_alpha=0, reg_lambda=1, ...)
```

Let's save our work before continuing.

```
jovian.commit()
```

QUESTION 8: Make predictions and evaluate your final model. If you're unhappy with the results, modify the hyperparameters above and try again.

```
rf2_train_preds = rf2.predict(train_inputs)
```

```
xg_train_preds = xg.predict(train_inputs)
```

```
rf2_train_rmse = mean_squared_error(rf2_train_preds, train_targets, squared=False)
```

```
xg_train_rmse = mean_squared_error(xg_train_preds, train_targets, squared=False)
```

```
rf2_val_preds = rf2.predict(val_inputs)
```

```
xg_val_preds = xg.predict(val_inputs)
```

```
rf2_val_rmse = mean_squared_error(rf2_val_preds, val_targets, squared=False)
```

```
xg_val_rmse = mean_squared_error(xg_val_preds, val_targets, squared=False)
```

```
print('Train RMSE: {}, Validation RMSE: {}'.format(rf2_train_rmse, rf2_val_rmse))
```

```
print('Train RMSE: {}, Validation RMSE: {}'.format(xg_train_rmse, xg_val_rmse))
```

Let's also view and plot the feature importances.

```
rf2_importance_df = pd.DataFrame({  
    'feature': train_inputs.columns,  
    'importance': rf2.feature_importances_  
}).sort_values('importance', ascending=False)
```


```
rf2_importance_df
```

	feature	importance
3	OverallQual	0.423416
15	GrLivArea	0.122802
25	GarageCars	0.074826
172	ExterQual_TA	0.042507
11	TotalBsmtSF	0.035542
...
159	Exterior2nd_Stone	0.000000
157	Exterior2nd_Other	0.000000
155	Exterior2nd_ImStucc	0.000000
149	Exterior2nd_AsphShn	0.000000
152	Exterior2nd_CBlock	0.000000

304 rows × 2 columns

```
sns.barplot(data=rf2_importance_df, x='importance', y='feature')
```

```
<AxesSubplot:xlabel='importance', ylabel='feature'>
```



Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on <https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

Make a Submission

To make a submission, just execute the following cell:

```
jovian.submit('zerotogbms-a2')
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on <https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

[jovian] Submitting assignment..

[jovian] Verify your submission at <https://jovian.ai/learn/machine-learning-with-python-zero-to-gbms/assignment/assignment-2-decision-trees-and-random-forests>

You can also submit your Jovian notebook link on the assignment page: <https://jovian.ai/learn/machine-learning-with-python-zero-to-gbms/assignment/assignment-2-decision-trees-and-random-forests>

Make sure to review the evaluation criteria carefully. You can make any number of submissions, and only your final submission will be evaluated.

Ask questions, discuss ideas and get help here: <https://jovian.ai/forum/c/zero-to-gbms/gbms-assignment-2/99>

NOTE: The rest of this assignment is optional.

Making Predictions on the Test Set

Let's make predictions on the test set provided with the data.

```
test_df = pd.read_csv('house-prices/test.csv')
```

test_df

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	Inside
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	Corner
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	Inside
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl	AllPub	Inside
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS	AllPub	Inside
...
1454	2915	160	RM	21.0	1936	Pave	NaN	Reg	Lvl	AllPub	Inside
1455	2916	160	RM	21.0	1894	Pave	NaN	Reg	Lvl	AllPub	Inside
1456	2917	20	RL	160.0	20000	Pave	NaN	Reg	Lvl	AllPub	Inside
1457	2918	85	RL	62.0	10441	Pave	NaN	Reg	Lvl	AllPub	Inside
1458	2919	60	RL	74.0	9627	Pave	NaN	Reg	Lvl	AllPub	Inside

1459 rows × 80 columns

First, we need to reapply all the preprocessing steps.

```
test_df[numeric_cols] = imputer.transform(test_df[numeric_cols])
test_df[numeric_cols] = scaler.transform(test_df[numeric_cols])
test_df[encoded_cols] = encoder.transform(test_df[categorical_cols])
```

/opt/conda/lib/python3.9/site-packages/pandas/core/frame.py:3678: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use `newframe = frame.copy()`

```
self[col] = igetitem(value, i)
```

```
test_inputs = test_df[numeric_cols + encoded_cols]
test_inputs
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtF
0	0.000000	0.202055	0.048246	0.444444	0.625	0.644928	0.183333	0.00000	0.08
1	0.000000	0.205479	0.060609	0.555556	0.625	0.623188	0.133333	0.06750	0.16
2	0.235294	0.181507	0.058566	0.444444	0.500	0.905797	0.800000	0.00000	0.14
3	0.235294	0.195205	0.040562	0.555556	0.625	0.913043	0.800000	0.01250	0.10

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtF
4	0.588235	0.075342	0.017318	0.777778	0.500	0.869565	0.700000	0.00000	0.04
...
1454	0.823529	0.000000	0.002973	0.333333	0.750	0.710145	0.333333	0.00000	0.00
1455	0.823529	0.000000	0.002776	0.333333	0.500	0.710145	0.333333	0.00000	0.04
1456	0.000000	0.476027	0.087406	0.444444	0.750	0.637681	0.766667	0.00000	0.21
1457	0.382353	0.140411	0.042726	0.444444	0.500	0.869565	0.700000	0.00000	0.05
1458	0.235294	0.181507	0.038921	0.666667	0.500	0.876812	0.733333	0.05875	0.13

1459 rows × 304 columns

We can now make predictions using our final model.

```
test_preds = rf2.predict(test_inputs)
```

```
xgtest_pred = xg.predict(test_inputs)
```

```
submission_df = pd.read_csv('house-prices/sample_submission.csv')
```

```
submission_df
```

	Id	SalePrice
0	1461	169277.052498
1	1462	187758.393989
2	1463	183583.683570
3	1464	179317.477511
4	1465	150730.079977
...
1454	2915	167081.220949
1455	2916	164788.778231
1456	2917	219222.423400
1457	2918	184924.279659
1458	2919	187741.866657

1459 rows × 2 columns

Let's replace the values of the `SalePrice` column with our predictions.

```
submission_df['SalePrice'] = xgtest_pred
```

Let's save it as a CSV file and download it.

```
submission_df.to_csv('submission.csv', index=False)
```

```
from IPython.display import FileLink
FileLink('submission.csv') # Doesn't work on Colab, use the file browser instead to download
submission.csv
```

We can now submit this file to the competition: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/submissions>

Name	Submitted	Wait time	Execution time	Score
submission.csv	just now	1 seconds	0 seconds	0.14842
Complete				

(OPTIONAL) QUESTION: Submit your predictions to the competition. Experiment with different models, feature engineering strategies and hyperparameters and try to reach the top 10% on the leaderboard.

Let's save our work before continuing.

```
jovian.commit()
```

[jovian] Updating notebook "ch20b112/python-random-forests-assignment" on <https://jovian.ai>

[jovian] Committed successfully! <https://jovian.ai/ch20b112/python-random-forests-assignment>

'<https://jovian.ai/ch20b112/python-random-forests-assignment>'

Making Predictions on Single Inputs

```
def predict_input(model, single_input):
    input_df = pd.DataFrame([single_input])
    input_df[numeric_cols] = imputer.transform(input_df[numeric_cols])
    input_df[numeric_cols] = scaler.transform(input_df[numeric_cols])
    input_df[encoded_cols] = encoder.transform(input_df[categorical_cols].values)
    return model.predict(input_df[numeric_cols + encoded_cols])[0]
```

```
sample_input = { 'MSSubClass': 20, 'MSZoning': 'RL', 'LotFrontage': 77.0, 'LotArea': 93
'Street': 'Pave', 'Alley': None, 'LotShape': 'IR1', 'LandContour': 'Lvl', 'Utilities':
'LotConfig': 'Inside', 'LandSlope': 'Gtl', 'Neighborhood': 'Names', 'Condition1': 'Nor
'BldgType': '1Fam', 'HouseStyle': '1Story', 'OverallQual': 4, 'OverallCond': 5, 'YearB
'YearRemodAdd': 1959, 'RoofStyle': 'Gable', 'RoofMatl': 'CompShg', 'Exterior1st': 'Ply
'Exterior2nd': 'Plywood', 'MasVnrType': 'None', 'MasVnrArea': 0.0, 'ExterQual': 'TA', 'Ex
```



```
'Foundation': 'CBlock', 'BsmtQual': 'TA', 'BsmtCond': 'TA', 'BsmtExposure': 'No', 'BsmtFinSF1': 569, 'BsmtFinType2': 'Unf', 'BsmtFinSF2': 0, 'BsmtUnfSF': 381, 'TotalBsmtSF': 950, 'Heating': 'GasA', 'HeatingQC': 'Fa', 'CentralAir': 'Y', 'Electrical': 'SBrkr', '2ndFlrSF': 0, 'LowQualFinSF': 0, 'GrLivArea': 1225, 'BsmtFullBath': 1, 'BsmtHalfBath': 1, 'BedroomAbvGr': 3, 'KitchenAbvGr': 1, 'KitchenQual': 'TA', 'TotRmsAbvGrd': 6, 'Fireplaces': 0, 'FireplaceQu': np.nan, 'GarageType': np.nan, 'GarageYrBlt': np.nan, 'GarageArea': 0, 'GarageQual': np.nan, 'GarageCond': np.nan, 'PavedDrive': 'Y', 'WoodDeckSF': 150, 'EnclosedPorch': 0, '3SsnPorch': 0, 'ScreenPorch': 0, 'PoolArea': 0, 'PoolQC': np.nan, 'MiscVal': 400, 'MoSold': 1, 'YrSold': 2010, 'SaleType': 'WD', 'SaleCondition': 'Normal'
```

```
predicted_price = predict_input(rf2, sample_input)
```

/opt/conda/lib/python3.9/site-packages/sklearn/base.py:450: UserWarning: X does not have valid feature names, but OneHotEncoder was fitted with feature names

```
warnings.warn(
```

/opt/conda/lib/python3.9/site-packages/pandas/core/frame.py:3678: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling `frame.insert` many times, which has poor performance. Consider joining all columns at once using `pd.concat(axis=1)` instead. To get a de-fragmented frame, use `newframe = frame.copy()`

```
self[col] = igetitem(value, i)
```

```
print('The predicted sale price of the house is ${}'.format(predicted_price))
```

The predicted sale price of the house is \$124455.83077923351

EXERCISE: Change the sample input above and make predictions. Try different examples and try to figure out which columns have a big impact on the sale price. Hint: Look at the feature importance to decide which columns to try.

Saving the Model

```
import joblib
```

```
house_prices_rf = {
    'model': rf2,
    'imputer': imputer,
    'scaler': scaler,
    'encoder': encoder,
    'input_cols': input_cols,
    'target_col': target_col,
    'numeric_cols': numeric_cols,
```

```
'categorical_cols': categorical_cols,  
'encoded_cols': encoded_cols  
}
```

```
joblib.dump(house_prices_rf, 'house_prices_rf.joblib')  
  
['house_prices_rf.joblib']
```

Let's save our work before continuing.

```
jovian.commit(outputs=['house_prices_rf.joblib'])
```

Predicting the Logarithm of Sale Price

(OPTIONAL) QUESTION: In the [original Kaggle competition](#), the model is evaluated by computing the Root Mean Squared Error on the logarithm of the sale price. Try training a random forest to predict the logarithm of the sale price, instead of the actual sales price and see if the results you obtain are better than the models trained above.